

Deep Learning Midterm Project Report

Che Hung Liao
cl8172@nyu.edu

Eleftherios Komvopoulos
ek4538@nyu.edu

Abstract

This report presents our approach to a deep learning task using modern neural network methods. We focus on building a competitive model under practical training-time constraints, compare multiple configurations, and summarize the design and optimization choices that contributed most to performance.

1 Introduction

This project focuses on solving a deep learning task using modern neural network approaches. The goal is to build a model that achieves strong performance on the given dataset while exploring different architectural and training choices.

Due to the computational cost of training, we emphasize iterative experimentation and practical improvements over exhaustive search. We evaluate different configurations and analyze what contributes most to performance.¹²³

2 Dataset

[TO BE FILLED]

We describe the dataset, preprocessing steps, and any augmentation techniques applied.

3 Model Description

We build on the Qwen2.5 series as our base model family, which provides strong general-purpose capabilities and a range of parameter sizes suitable for experimentation on limited hardware. Our experiments span both general-purpose instruction-tuned models and code-specialized instruction-tuned models within the Qwen2.5 family. In particular, Qwen2.5-1.5B-Instruct serves as a general instruction-following baseline, whereas Qwen2.5-Coder-3B-Instruct is more specialized for code-like

and syntax-constrained generation. This distinction is important for our task, since SVG generation behaves more like structured markup generation than free-form natural language generation: the model must maintain valid tags, attributes, nesting, and quotation syntax while still following the semantic content of the prompt. For this reason, the coder-oriented variants are a natural fit for later-stage experiments that emphasize structured SVG output. We initially evaluated 1.5B, 3B, and 7B parameter models; however, the 7B version consistently pushed GPU memory to near out-of-memory limits, so we ultimately standardized on 3B and 1.5B models as the best balance between capacity and stable training/inference.

4 Experimentation

Our experimentation followed a staged and hypothesis-driven progression. We began with the official starter implementation and first modified it so that it could run reliably on a single AMD MI100 GPU, since the original workflow was not directly aligned with our hardware environment. Using this adapted starter as the baseline, we conducted a sequence of controlled ablations on the Qwen2.5-1.5B model family. We first increased the maximum sequence length from 1536 to 2048 to test whether longer context windows would improve SVG generation quality. We then increased the LoRA rank and LoRA scaling factor to evaluate whether stronger low-rank adaptation would improve performance. Finally, we expanded the training set size in stages, moving from the default subset to 24,000 and then 36,000 samples, in order to determine whether additional supervision translated into better results. The full set of experiment configurations is summarized in Table 1. These initial experiments established a clear reference point, but they did not yield a meaningful improvement over the starter baseline. In these starter experi-

¹Code repository: [GitHub repository](#).

²Model weights: [GitHub adapter link](#).

³We used GPT and Cursor for implementation, debugging, learning, report drafting, and proofreading.

ments, we set the evaluation batch size to 1 in order to minimize the risk of out-of-memory failures during validation, which we observed frequently when larger evaluation batches were used on the AMD platform.

We therefore shifted to Kaggle’s NVIDIA T4 $\times$ 2 environment and moved to the Unsloth 3B model family. In this phase, we used 4-bit loading so that the larger 3B models would remain trainable within the memory limits of the hosted GPUs while reducing the risk of out-of-memory failures. Starting from an initial 3B configuration, we increased LoRA rank and per-device batch size to better utilize the available GPU resources and improve training throughput. We then introduced a more explicit system prompt containing SVG-specific rules to test whether stronger task guidance improved both training behavior and inference quality. In a later Kaggle variant, we further increased the per-device batch size, doubled the learning rate, and doubled the number of training epochs in order to test a more aggressive optimization setup while remaining within the platform’s memory constraints.

As the project progressed, Kaggle usage limits became a practical bottleneck for longer and more resource-intensive runs. We therefore returned to the AMD platform, this time using a two-GPU setup for our final high-capacity experiment. In this last stage, we increased the maximum sequence length to 3072 and switched to Qwen2.5-Coder-3B-Instruct, while also lowering the LoRA rank, increasing the LoRA scaling factor, and training for more epochs. This configuration was intended to combine longer context, stronger code-oriented reasoning, and a more stable training environment for extended runs. The switch to the coder-specialized model was motivated by the observation that SVG generation is fundamentally a structured markup task, so a model with a stronger bias toward code-like syntax and constrained output was a better fit than a general instruction-tuned alternative. We again set the evaluation batch size to 1 in this final AMD-based configuration for the same reason: to reduce evaluation-time memory pressure and avoid out-of-memory interruptions.

5 Results

The results in Table 1 show a clear progression across the different stages of experimentation. The starter-based AMD runs were stable but did not show measurable improvement across changes in

sequence length, LoRA configuration, or training-set size: all five starter variants produced the same private score of 8.42524 and public score of 11.11991. This suggests that, within that baseline setup, simply scaling context length, adapter strength, or sample count was not sufficient to improve performance.

Moving to the Kaggle NVIDIA T4 $\times$ 2 environment with the 3B Instruct model produced an immediate improvement over the starter baseline. The initial 3B configuration achieved a private score of 7.27635 and a public score of 9.65655. Increasing LoRA rank and batch size further improved performance to 8.27318 private and 10.84663 public. The strongest Kaggle result came from the prompt-augmented variant, which reached 9.79340 private and 12.23749 public, indicating that explicit SVG-oriented instruction formatting contributed meaningfully to model behavior. By contrast, the later variant with larger batch size, higher learning rate, and more epochs did not improve further, ending at 9.68681 private and 11.76432 public, which suggests that more aggressive optimization alone did not yield additional gains.

The best overall performance came from the final AMD-based experiment using Qwen2.5-Coder-3B-Instruct with a 3072-token context window. This configuration achieved 10.77991 on the private leaderboard and 13.11993 on the public leaderboard, outperforming all previous runs. Taken together, these results indicate that the most important improvements came not from isolated scaling of baseline hyperparameters, but from a combination of stronger model choice, task-aligned prompting, and a final transition to a coder-specialized model with longer context and a higher-capacity training setup.

6 Conclusion

In this project, we explored different modeling approaches and training strategies to improve performance on the task.

We found that careful tuning of hyperparameters and iterative experimentation were key to achieving better results. However, training time posed a significant constraint, limiting the number of experiments that could be conducted.

Some approaches did not lead to improvements, highlighting the importance of targeted experimentation rather than blindly increasing complexity.

In future work, we would implement better ex-

Experiment	Model	Max Seq	r	α	LR	Epochs	Train BS	Eval BS	Grad Accum	Target Samples	Score (Private)	Score (Public)	Platform
starter/1536	Qwen2.5-1.5B-Instruct	1536	16	16	2e-4	1	1	1	8	12000	8.42524	11.11991	AMD MI100
starter/2048	Qwen2.5-1.5B-Instruct	2048	16	16	2e-4	1	1	1	8	12000	8.42524	11.11991	AMD MI100
starter/2048-r32	Qwen2.5-1.5B-Instruct	2048	32	64	2e-4	1	1	1	8	12000	8.42524	11.11991	AMD MI100
starter/2048-r32-24000	Qwen2.5-1.5B-Instruct	2048	32	64	2e-4	1	1	1	8	24000	8.42524	11.11991	AMD MI100
starter/2048-r32-36000	Qwen2.5-1.5B-Instruct	2048	32	64	2e-4	1	1	1	8	36000	8.42524	11.11991	AMD MI100
2gpu version	Qwen2.5-3B-Instruct-bnb-4bit	2048	16	16	2e-4	1	1	8 (default)	4	12000	7.27635	9.65655	Kaggle NVIDIA T4 $\times$ 2
2gpu-version-batch2	Qwen2.5-3B-Instruct-bnb-4bit	2048	32	16	2e-4	1	2	8 (default)	4	12000	8.27318	10.84663	Kaggle NVIDIA T4 $\times$ 2
2gpu-version-batch2-with-prompt	Qwen2.5-3B-Instruct-bnb-4bit	2048	32	16	2e-4	1	2	8 (default)	4	12000	9.79340	12.23749	Kaggle NVIDIA T4 $\times$ 2
2gpu-version-batch32-2xlr-step4	Qwen2.5-3B-Instruct-bnb-4bit	2048	32	16	4e-4	2	4	8 (default)	4	12000	9.68681	11.76432	Kaggle NVIDIA T4 $\times$ 2
3072-17hrs	Qwen2.5-Coder-3B-Instruct	3072	16	64	2e-4	3	4	1	4	20000	10.77991	13.11993	AMD MI100 $\times$ 2

Table 1: Summary of experiment configurations, training platforms, and Kaggle competition submission scores. Here, r denotes LoRA rank, α denotes LoRA scaling factor, LR denotes learning rate, Train BS denotes per-device training batch size, Eval BS denotes per-device evaluation batch size, Grad Accum denotes gradient accumulation steps, and Target Samples denotes the number of training examples used in that experiment. Score (Private) and Score (Public) refer to the private and public leaderboard scores reported by Kaggle for each submission.

periment tracking and explore more efficient training methods to allow broader exploration of the design space.